

1 A Additional Details on Modeling

2 In this section, we provide additional details on our modeling techniques.

3 A.1 Soft Nearest Neighbor Loss

4 The soft nearest neighbor loss applied to the learned encodings z_j^i looks like:

$$l_{snn} = -\frac{1}{N} \sum_{i,j} \log \left(\frac{\sum_{k \neq i} e^{-\frac{\|z_j^i - z_j^k\|^2}{T}}}{\sum_{\substack{k,l \\ (k,l) \neq (i,j)}} e^{-\frac{\|z_j^i - z_l^k\|^2}{T}}} \right)$$

5 where T is a temperature parameter.

6 In our experiments, we combine this loss function with the reconstruction loss from the autoencoder
7 so that the final loss becomes:

$$\mathcal{J} = \frac{1}{N} \sum_{i,j} \left(\|\tilde{x}_j^i - x_j^i\|^2 - \lambda \log \left(\frac{\sum_{k \neq i} e^{-\frac{\|z_j^i - z_j^k\|^2}{T}}}{\sum_{\substack{k,l \\ (k,l) \neq (i,j)}} e^{-\frac{\|z_j^i - z_l^k\|^2}{T}}} \right) \right)$$

8 where λ is a regularization coefficient.

9 A.2 Hybrid Architectures

10 Finally, we propose a hybrid architecture between siamese neural networks and autoencoders. The
11 idea is to augment the siamese neural network by adding a decoder so that we can have add a
12 reconstruction term to the loss function. We also add the other terms related to our customized
13 disentanglement. The complete loss function looks like:

$$\mathcal{J} = \sum_{a=1}^n \sum_{i=1}^{I_s} \sum_{j \neq a} L_T(x_a^i, x_a^k, x_j^i) + \gamma L_R(x_a^i, x_a^k, x_j^i) + \lambda L_C(v_a^i, v_a^k, v_j^i)$$

14 with L_T as provided in Section 4.4, L_R is the reconstruction of the 3 terms (anchor, positive, and
15 negative) and L_C corresponds to the configuration approximation for the 3 terms as well.

$$L_R(x_a^i, x_a^k, x_j^i) = \|\tilde{x}_a^i - x_a^i\|^2 + \|\tilde{x}_a^k - x_a^k\|^2 + \|\tilde{x}_j^i - x_j^i\|^2$$

$$L_C(v_a^i, v_a^k, v_j^i) = \|\tilde{v}_a^i - v_a^i\|^2 + \|\tilde{v}_a^k - v_a^k\|^2 + \|\tilde{v}_j^i - v_j^i\|^2$$

16 A.3 Regression on end target objective

17 For the encoder/decoder based architectures (namely the auto-encoders, siamese neural networks
18 and hybrid aproaches) we have introduced in earlier sections losses that correspond to learning the
19 representations. The full architecture however has a regression module that takes learned workload
20 encodings z_j^i and fits a regression function on runtime latency. The regressor takes as input the job
21 configuration v_j^i and z_j (the centroid of $\{z_j^i\}_i$ for a particular workload j) and tries to approximate at
22 its output the runtime latency y_j^i .

23 The loss function for the regression is:

$$L = \frac{1}{N} \sum_{i,j} (f(v_j^i, z_j) - y_j^i)^2$$

24 With our best performing representation learning technique (siamese neural network), we tried both:
25 (1) training separately (a) the siamese neural network and (b) the regression architecture (2) training
26 the encoder first, then finetuning its layers when minimizing the loss function of the regressor. Fine
27 tuning the encoder while training the regressor didn't give improvements on the test set errors for the
28 end regression task.

29 B Additional Experimental Details

30 B.1 More Details regarding the traces

31 All of the traces have been collected from Spark clusters deployed on homogeneous hardware. Each
32 Spark cluster spans 3 nodes (with 1 node reserved for the driver and 2 nodes left for executors). The
33 traces contain two types of metrics: (1) Spark related metrics collected using Spark listener and (2)
34 OS related metrics collected using the unix command *nmon*.

35 **Preprocessing** The traces are first averaged across the period of execution of the running Spark
36 workloads. Then, we drop metrics for which the values are either constant or NaN across different
37 traces. Then, we do minmaxscaling of both the knobs v_j^i as well as the retained runtime metrics x_j^i .

38 Within the streaming trace, we vary 10 knobs, and thus the dimension of v_j^i is 10. The total number
39 of metrics in raw traces is 931 metrics. This number becomes 561 after preprocessing. Thus, the
40 dimension of x_j^i is 561.

41 On the other hand, within the TPCx-BB trace we vary 12 knobs. Thus, the dimension of v_j^i is 12.
42 The total number of metrics in raw traces is however 572 metrics. This number becomes 286 after
43 applying the preprocessing. Thus, the dimension of x_j^i is 286.

44 B.2 Hyper-parameter tuning

45 For the encoder/decoder based architectures as well as the neural networks we tuned:

- 46 (1) topology related hyper-parameters (number of layers, number of hidden units per layer, activations)
- 47 (2) learning related hyper-parameters (learning rate, number of epochs, patience (for early stopping))
- 48 (3) loss related hyper-parameters (regularization coefficients introduced in losses having more than
- 49 one term, temperature, etc...)

50 We tuned the hyper-parameters by random sampling from a pool of hyper-parameters. We use a
51 5-fold cross validation scheme for tuning our workloads. We try to simulate the same conditions
52 on the test set when we do cross validation, thus we consider observing only few (1 or 5 traces) for
53 training workloads within the left out fold during the cross validation procedure.

54 It may appear that we are casting tuning Spark workloads to tuning hyper-parameters of machine
55 learning models. While this is partially correct, it is important to emphasize that the machine learning
56 solution to modeling performances of spark workloads we are proposing can be tuned overnight (and
57 not at the time of the execution of the Spark workload). Having a robust global model that allows us
58 to predict performances of a new submitted Spark job from a unique (or few) trace, makes the tuning
59 of Spark workload non-invasive to the user, and much faster.

60 C Implementation and Hardware Details

61 We have implemented all encoder/decoder based architectures as well as the neural network regressor
62 using Tensorflow [1]. For the comparison with baseline representation learning techniques (PCA and
63 KPCA), we used open source implementations from scikit-learn [16].

64 We have a dedicated server of 20 nodes for training our models. Each node has 2 x Intel(R) Xeon(R)
65 Gold 6130 CPU @ 2.10GHz processors (with 16 cores per processor) and 754 GB of RAM.